

Practical Artificial Intelligence of Image Recognition, Spring 2023 HW4

National Taiwan University B11705009 梁安哲

```
In [ ]: # import modules
        from keras.datasets import cifar10
        import numpy as np
        import tensorflow as tf
        import matplotlib.pyplot as plt
```

```
In [ ]: # prepare dataset & labels
        from keras.utils import np_utils

        (x_img_train,y_label_train),(x_img_test,y_label_test)=cifar10.load_data()
        x_img_train_normalize = x_img_train.astype('float32') / 255.0
        x_img_test_normalize = x_img_test.astype('float32') / 255.0
        y_label_train_OneHot = np_utils.to_categorical(y_label_train)
        y_label_test_OneHot = np_utils.to_categorical(y_label_test)
```

Part 1: MLP model

```
In [ ]: # construct the MLP model with 3 128 node fully connected layers
        model = tf.keras.models.Sequential([
            tf.keras.layers.Flatten(),
            tf.keras.layers.Dense(128, activation='relu'),
            tf.keras.layers.Dense(128, activation='relu'),
            tf.keras.layers.Dense(128, activation='relu'),
            tf.keras.layers.Dropout(0.2),
            tf.keras.layers.Dense(10, activation="softmax")
        ])
```

```
In [ ]: # compile the model with loss function: categorical_crossentropy and optim
        model.compile(optimizer='adam',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])
```

```
In [ ]: # train the model with epoch:10 batch_size: 128
        train_history = history = model.fit(x_img_train_normalize, y_label_train_OneHot,
                                             validation_split=0.2,
                                             epochs=10, batch_size=128, verbose=1)
```

```

Epoch 1/10
313/313 [=====] - 3s 7ms/step - loss: 1.9299 - ac
curacy: 0.2942 - val_loss: 1.8035 - val_accuracy: 0.3497
Epoch 2/10
313/313 [=====] - 2s 6ms/step - loss: 1.7451 - ac
curacy: 0.3728 - val_loss: 1.7160 - val_accuracy: 0.3885
Epoch 3/10
313/313 [=====] - 2s 6ms/step - loss: 1.6610 - ac
curacy: 0.4045 - val_loss: 1.6403 - val_accuracy: 0.4213
Epoch 4/10
313/313 [=====] - 2s 6ms/step - loss: 1.6057 - ac
curacy: 0.4252 - val_loss: 1.5778 - val_accuracy: 0.4393
Epoch 5/10
313/313 [=====] - 2s 6ms/step - loss: 1.5617 - ac
curacy: 0.4401 - val_loss: 1.5542 - val_accuracy: 0.4437
Epoch 6/10
313/313 [=====] - 1s 5ms/step - loss: 1.5365 - ac
curacy: 0.4473 - val_loss: 1.5658 - val_accuracy: 0.4380
Epoch 7/10
313/313 [=====] - 1s 4ms/step - loss: 1.5148 - ac
curacy: 0.4578 - val_loss: 1.5317 - val_accuracy: 0.4571
Epoch 8/10
313/313 [=====] - 2s 6ms/step - loss: 1.4848 - ac
curacy: 0.4699 - val_loss: 1.6246 - val_accuracy: 0.4293
Epoch 9/10
313/313 [=====] - 2s 6ms/step - loss: 1.4688 - ac
curacy: 0.4735 - val_loss: 1.5838 - val_accuracy: 0.4344
Epoch 10/10
313/313 [=====] - 2s 6ms/step - loss: 1.4516 - ac
curacy: 0.4799 - val_loss: 1.5271 - val_accuracy: 0.4574

```

```

In [ ]: import matplotlib.pyplot as plt
        # functions for data visualization
        def show_acc_train_history(train_acc,test_acc):
            plt.plot(train_history.history[train_acc])
            plt.plot(train_history.history[test_acc])
            plt.title('Train History')
            plt.ylabel('Accuracy')
            plt.xlabel('Epoch')
            plt.legend(['train', 'test'], loc='upper left')
            plt.show()

```

```


In [ ]: def show_loss_train_history(train_loss,test_loss):
            plt.plot(train_history.history[train_loss])
            plt.plot(train_history.history[test_loss])
            plt.title('Train History')
            plt.ylabel('Loss')
            plt.xlabel('Epoch')
            plt.legend(['train', 'test'], loc='upper left')
            plt.show()

```


```

In [ ]: show_acc_train_history('accuracy','val_accuracy')

```



```
In [ ]: show_loss_train_history('loss', 'val_loss')
```



Part 2: CNN model

```
In [ ]: data_augmentation = tf.keras.Sequential([
    tf.keras.layers.RandomFlip("horizontal_and_vertical"),
    tf.keras.layers.RandomRotation(0.1),
])
# construct the CNN model with data_augmentation
model = tf.keras.models.Sequential([
    data_augmentation,
    tf.keras.layers.Conv2D(filters=32, kernel_size=(3,3),
        input_shape=(32, 32,3),
        activation='relu',
        padding='same'),
    tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),
    tf.keras.layers.Conv2D(filters=64, kernel_size=(3, 3),
        activation='relu', padding='same'),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(2048, activation='relu'),
    tf.keras.layers.Dense(1024, activation='relu'),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10 , activation="softmax")
])
```

```
In [ ]: # compile the model with loss funcion: categorical_crossentropy and optim
model.compile(optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy'])
```

```
In [ ]: # train the model with epoch:10 batch_size: 64
train_history = model.fit(x_img_train_normalize, y_label_train_OneHot,
    validation_split=0.2,
    epochs=10, batch_size=64, verbose=1)
```

```
Epoch 1/10
625/625 [=====] - 12s 15ms/step - loss: 1.7878 - accuracy: 0.3409 - val_loss: 1.5190 - val_accuracy: 0.4411
Epoch 2/10
625/625 [=====] - 9s 15ms/step - loss: 1.5063 - accuracy: 0.4514 - val_loss: 1.3935 - val_accuracy: 0.4939
Epoch 3/10
625/625 [=====] - 9s 15ms/step - loss: 1.4001 - accuracy: 0.4981 - val_loss: 1.3142 - val_accuracy: 0.5297
Epoch 4/10
625/625 [=====] - 9s 15ms/step - loss: 1.3166 - accuracy: 0.5293 - val_loss: 1.2838 - val_accuracy: 0.5464
Epoch 5/10
625/625 [=====] - 9s 15ms/step - loss: 1.2551 - accuracy: 0.5510 - val_loss: 1.1902 - val_accuracy: 0.5847
Epoch 6/10
625/625 [=====] - 9s 15ms/step - loss: 1.2067 - accuracy: 0.5685 - val_loss: 1.2346 - val_accuracy: 0.5654
Epoch 7/10
625/625 [=====] - 9s 15ms/step - loss: 1.1527 - accuracy: 0.5893 - val_loss: 1.1490 - val_accuracy: 0.5996
Epoch 8/10
625/625 [=====] - 9s 15ms/step - loss: 1.1216 - accuracy: 0.6015 - val_loss: 1.1350 - val_accuracy: 0.6029
Epoch 9/10
625/625 [=====] - 9s 15ms/step - loss: 1.0802 - accuracy: 0.6183 - val_loss: 1.0909 - val_accuracy: 0.6229
Epoch 10/10
625/625 [=====] - 9s 15ms/step - loss: 1.0479 - accuracy: 0.6281 - val_loss: 1.1610 - val_accuracy: 0.6092
```

```
In [ ]: model.summary()
```

Model: "sequential_8"


Layer (type)	Output Shape	Param #
sequential_7 (Sequential)	(None, 32, 32, 3)	0
conv2d_4 (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d_2 (MaxPooling 2D)	(None, 16, 16, 32)	0
conv2d_5 (Conv2D)	(None, 16, 16, 64)	18496
flatten_5 (Flatten)	(None, 16384)	0

Layer (type)	Output Shape	Param #
sequential_7 (Sequential)	(None, 32, 32, 3)	0
conv2d_4 (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d_2 (MaxPooling 2D)	(None, 16, 16, 32)	0
conv2d_5 (Conv2D)	(None, 16, 16, 64)	18496
flatten_5 (Flatten)	(None, 16384)	0
dense_20 (Dense)	(None, 2048)	33556480
dense_21 (Dense)	(None, 1024)	2098176
dense_22 (Dense)	(None, 512)	524800
dropout_5 (Dropout)	(None, 512)	0
dense_23 (Dense)	(None, 10)	5130


Total params: 36,203,978
Trainable params: 36,203,978
Non-trainable params: 0

This is the architecture of our model.

```
In [ ]: show_acc_train_history('accuracy', 'val_accuracy')
```



```
In [ ]: show_loss_train_history('loss', 'val_loss')
```



We can see that our CNN model outperform the MLP model (62% v.s. 48%), and from the accuracy graph, we can see that data augmentation makes our model less prone to

overfitting.